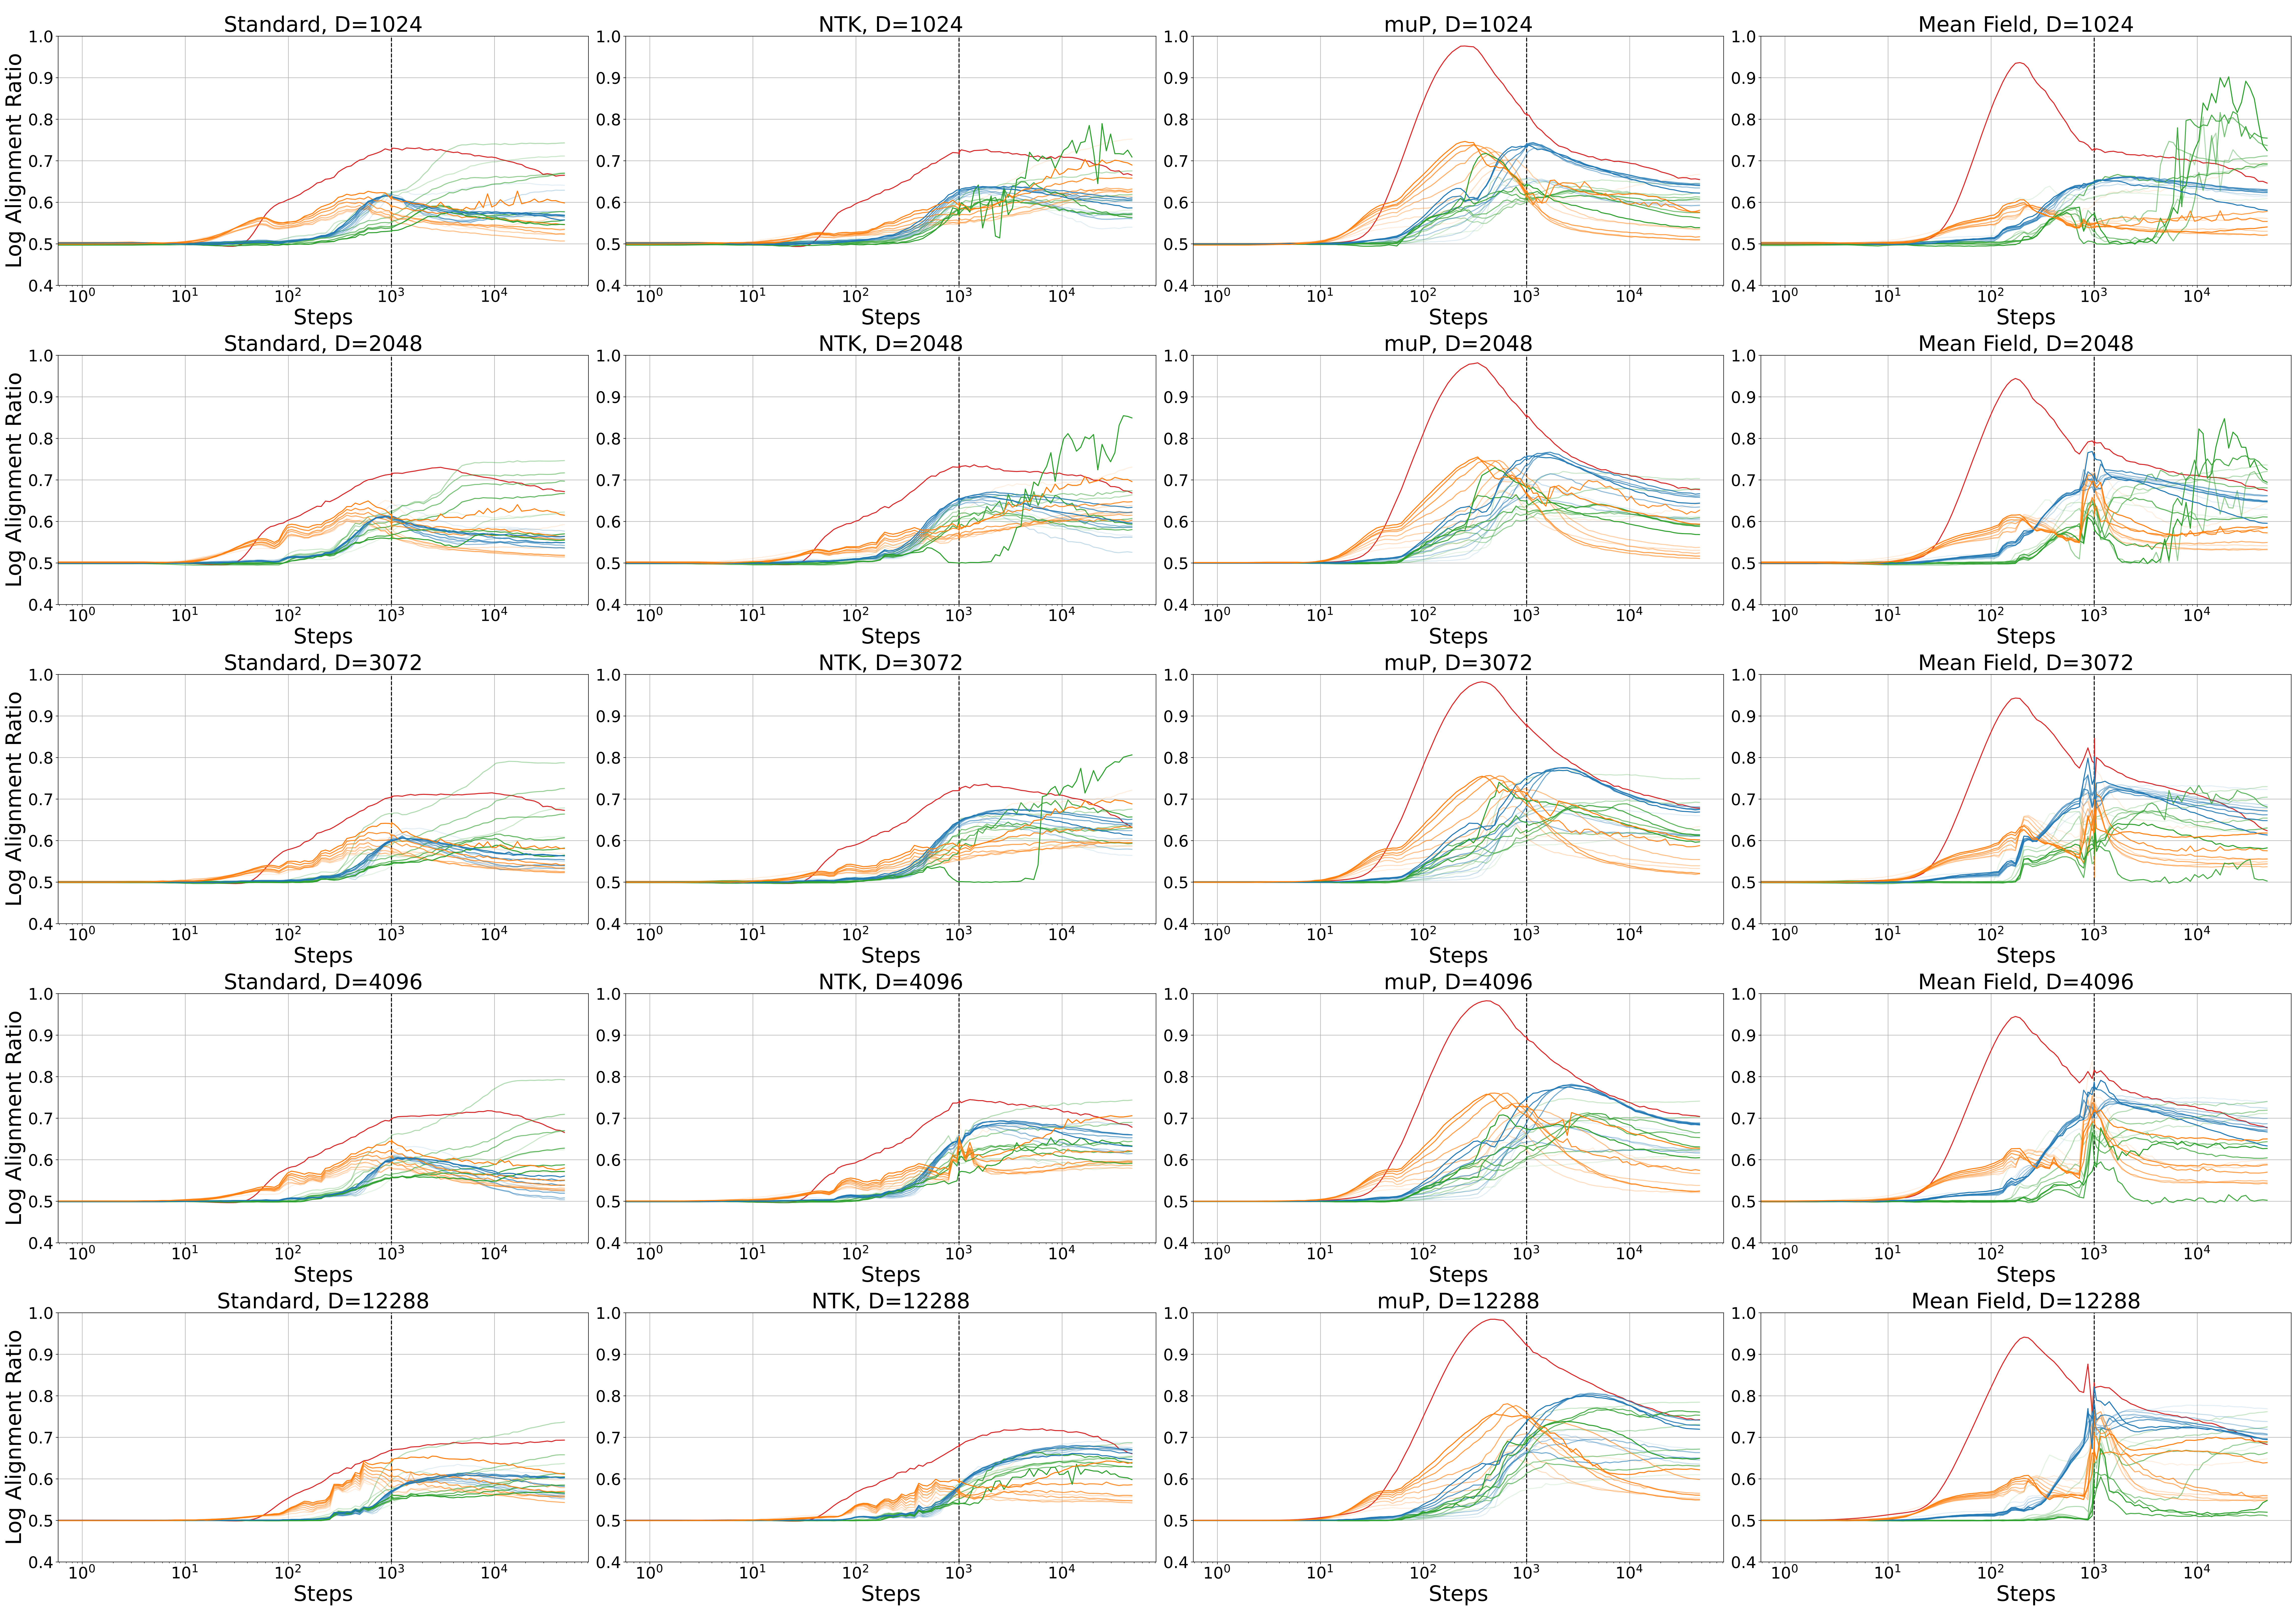


Adam Alignment Measurements (at optimal learning rates)



----- end of warmup B0/key B0/mlp_1 B1/mlp_0 B2/key B2/mlp_1 B3/mlp_0 B4/key B4/mlp_1 B5/mlp_0 B6/key B6/mlp_1 B7/mlp_0
 ——— readout B0/mlp_0 B1/key B1/mlp_1 B2/mlp_0 B3/key B3/mlp_1 B4/mlp_0 B5/key B5/mlp_1 B6/mlp_0 B7/key B7/mlp_1